

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df=pd.read_csv(r"C:\Users\venka\OneDrive\Documents\fiat500_VehicleSelection_D
print(df)
```

	ID	model	engine_power	age_in_days	km	previous_owners
0	1	lounge	51	882	25000	1 \
1	2	pop	51	1186	32500	1
2	3	sport	74	4658	142228	1
3	4	lounge	51	2739	160000	1
4	5	pop	73	3074	106880	1
...
1533	1534	sport	51	3712	115280	1
1534	1535	lounge	74	3835	112000	1
1535	1536	pop	51	2223	60457	1
1536	1537	lounge	51	2557	80750	1
1537	1538	pop	51	1766	54276	1

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

```
In [3]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing,svm
```

```
In [4]: df=df[['km','price']]
df.columns=['Km','Price']
```

```
In [5]: df.head(10)
```

```
Out[5]:
```

	Km	Price
0	25000	8900
1	32500	8800
2	142228	4200
3	160000	6000
4	106880	5700
5	70225	7900
6	11600	10750
7	49076	9190
8	76000	5600
9	89000	6000

```
In [6]: df.tail()
```

```
Out[6]:
```

	Km	Price
1533	115280	5200
1534	112000	4600
1535	60457	7500
1536	80750	5990
1537	54276	7900

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1538 entries, 0 to 1537  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0    Km      1538 non-null    int64  
1   Price   1538 non-null    int64  
dtypes: int64(2)  
memory usage: 24.2 KB
```

```
In [8]: df.describe()
```

```
Out[8]:
```

	Km	Price
count	1538.000000	1538.000000
mean	53396.011704	8576.003901
std	40046.830723	1939.958641
min	1232.000000	2500.000000
25%	20006.250000	7122.500000
50%	39031.000000	9000.000000
75%	79667.750000	10000.000000
max	235000.000000	11100.000000

```
In [9]: df.shape
```

```
Out[9]: (1538, 2)
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: Km      0
Price      0
dtype: int64
```


```
In [11]: x=np.array(df['Km']).reshape(-1,1)
y=np.array(df['Price']).reshape(-1,1)
```

```
In [12]: df.dropna(inplace=True)
```

```
In [13]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```


```
0.7522535193571669
```

```
In [14]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [15]: udf=df[:500]
sns.lmplot(x="Km",y="Price",data=udf,order=1,ci=None)
```

Out[15]: <seaborn.axisgrid.FacetGrid at 0x21141ef4110>




```
In [16]: udf.fillna(method='ffill',inplace=True)
X=np.array(udf['Km']).reshape(-1,1)
y=np.array(udf['Price']).reshape(-1,1)
udf.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3)
regr.fit(X_train,y_train)
```

Out[16]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.


```
In [17]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [18]: from sklearn.linear_model import Ridge,Lasso,RidgeCV,LassoCV
```

```
In [19]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[19]: <Axes: >



```
In [20]: from sklearn.preprocessing import StandardScaler
features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X_train is (1076, 2)

The dimension of X_test is (462, 2)

```
In [21]: #Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 1.0

The test score for Linear model is 1.0

```
In [22]: #ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.9997095924476731

The test score for ridge model is 0.9997198323998524


```
In [23]: #using the linear cv model for ridge regression
from sklearn.linear_model import RidgeCV
#ridge cross validation
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

0.9999999999999966


0.9999999999999674

```
In [24]: #using the linear cv model for lasso regression
from sklearn.linear_model import LassoCV
#Lasso cross validation
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)
#score
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

0.9999999877496772


0.9999999874481674

```
In [36]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',marker:
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize:
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



```
In [26]: #ridge regression  
plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[26]: <Axes: >




```
In [27]: #Lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.9999728562194999


The test score for lasso model is 0.9999728508562553

```
In [35]: plt.figure(figsize=(10,10))
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='o',marker:
plt.plot(features,reg.coef_,alpha=0.5,linestyle='none',marker='*',markersize:
plt.xticks(rotation=90)
plt.legend()
plt.show()
```

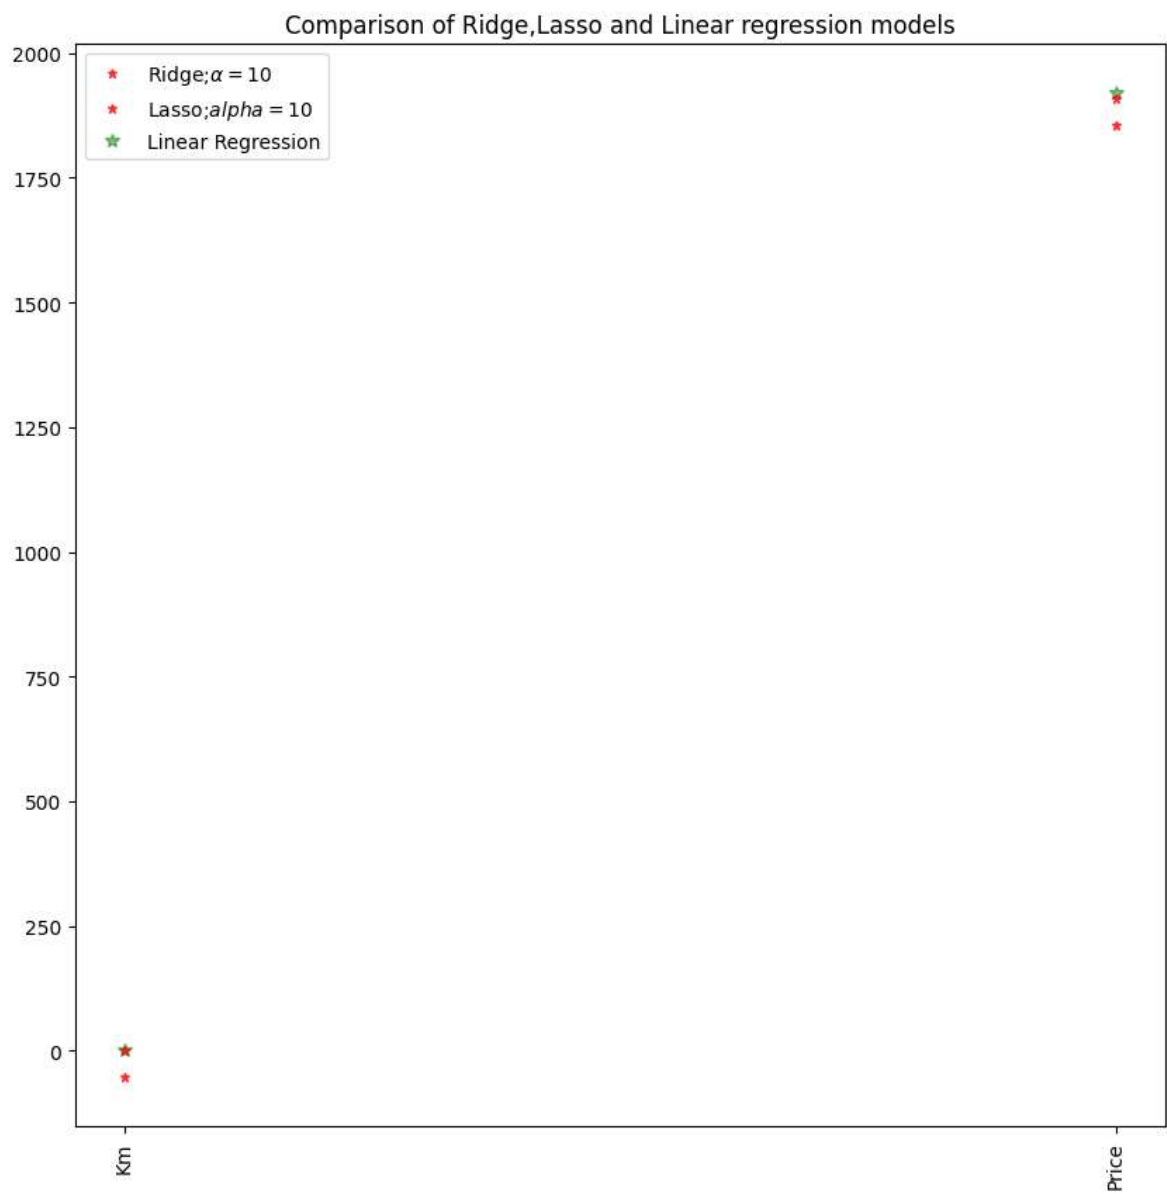


```
In [29]: pd.Series(lassoReg.coef_, features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[29]: <Axes: >
```



```
In [32]: #plot size
plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',marker:
#add plot for Lasso regression
plt.plot(features,lassoReg.coef_,alpha=0.7,linestyle='none',marker='*',marker:
#add plot for Linear model
plt.plot(features,regr.coef_,alpha=0.5,linestyle='none',marker='*',markersize:
#rotate axis
plt.xticks(rotation=90)
plt.legend()
plt.title("Comparison of Ridge,Lasso and Linear regression models")
plt.show()
```



In []: